

```
In [ ]: !pip install transformers -U
```

```
In [ ]: import pandas as pd
```

```
In [ ]: # https://www.kaggle.com/competitions/jigsaw-toxic-comment-classification-challenge
data = pd.read_csv("/content/train.csv", engine="python")
data.head()
```

```
In [ ]: data['toxic'].value_counts()
```

```
In [ ]: data = data[['comment_text', 'toxic']]
data = data[0:1000]
data.head()
```

```
In [ ]: import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, recall_score, precision_score, f1_score
import torch
from transformers import TrainingArguments, Trainer
from transformers import BertTokenizer, BertForSequenceClassification
```

```
In [ ]: from transformers import BertTokenizer, BertForSequenceClassification
tokenizer = BertTokenizer.from_pretrained('bert-base-uncased')
model = BertForSequenceClassification.from_pretrained('bert-base-uncased', num_label
```

```
In [ ]: model
```

```
In [ ]: # model = model.to('cuda')
```

```
In [ ]: sample_data = ["I am eating", "I am playing "]
tokenizer(sample_data, padding=True, truncation=True, max_length=512)
```

```
In [ ]: X = list(data["comment_text"])
y = list(data["toxic"])
X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.2, stratify=y)
X_train_tokenized = tokenizer(X_train, padding=True, truncation=True, max_length=51
X_val_tokenized = tokenizer(X_val, padding=True, truncation=True, max_length=512)
```

```
In [ ]: X_train_tokenized.keys()
```

```
In [ ]: print(X_train_tokenized['attention_mask'][0])
```

```
In [ ]: len(X_train), len(X_val)
```

```
In [ ]: # Create dataset
class Dataset(torch.utils.data.Dataset):
    def __init__(self, encodings, labels=None):
        self.encodings = encodings
        self.labels = labels
```

```

def __getitem__(self, idx):
    item = {key: torch.tensor(val[idx]) for key, val in self.encodings.items()}
    if self.labels:
        item["labels"] = torch.tensor(self.labels[idx])
    return item

def __len__(self):
    return len(self.encodings["input_ids"])

```

```

In [ ]: train_dataset = Dataset(X_train_tokenized, y_train)
        val_dataset = Dataset(X_val_tokenized, y_val)

```

```

In [ ]: train_dataset[5]

```

```

In [ ]: def compute_metrics(p):
        print(type(p))
        pred, labels = p
        pred = np.argmax(pred, axis=1)

        accuracy = accuracy_score(y_true=labels, y_pred=pred)
        recall = recall_score(y_true=labels, y_pred=pred)
        precision = precision_score(y_true=labels, y_pred=pred)
        f1 = f1_score(y_true=labels, y_pred=pred)

        return {"accuracy": accuracy, "precision": precision, "recall": recall, "f1": f

```

```

In [ ]: # Define Trainer
        args = TrainingArguments(
            output_dir="output",
            num_train_epochs=1,
            per_device_train_batch_size=8
        )
        trainer = Trainer(
            model=model,
            args=args,
            train_dataset=train_dataset,
            eval_dataset=val_dataset,
            compute_metrics=compute_metrics
        )

```

```

In [ ]: trainer.train()

```

```

In [ ]: trainer.evaluate()

```

```

In [ ]: np.set_printoptions(suppress=True)

```

```

In [ ]: text = "That was good point"
        # text = "go to hell"
        inputs = tokenizer(text, padding = True, truncation = True, return_tensors='pt').to(
        outputs = model(**inputs)
        print(outputs)
        predictions = torch.nn.functional.softmax(outputs.logits, dim=-1)
        print(predictions)

```

```
predictions = predictions.cpu().detach().numpy()
if predictions[0][0] > 0.5:
    print("Not Toxic")
else:
    print("Toxic")
```

```
In [ ]: trainer.save_model('CustomModel')
```

```
In [ ]: model_2 = BertForSequenceClassification.from_pretrained("CustomModel")
model_2.to('cuda')
```

```
In [ ]: # text = "That was good point"
inputs = tokenizer(text,padding = True, truncation = True, return_tensors='pt').to(
outputs = model_2(**inputs)
predictions = torch.nn.functional.softmax(outputs.logits, dim=-1)
predictions = predictions.cpu().detach().numpy()
if predictions[0][0] > 0.5:
    print("Not Toxic")
else:
    print("Toxic")
```

```
In [ ]:
```